

Causal Inference & ML

TA Session 2

Data Science for Decision Making

Lucia Sauer

2026-02-13

Outline

- Non Linear Models
- Nonparametric Density Estimation
 - Histograms
 - Kernell Density Estimation
- Non Parametric Regression Models
 - binscatter
 - NW estimator
 - Local Linear Regression
- Semiparametric Regression Models
 - Robinson's difference estimator

Non Linear Models

The Data

- **Dataset:** CRIME4 (Wooldridge)
- **Panel structure:** 90 counties, T=7 years (1981-1987)
- **Outcome variable:** Binary indicator for high crime rate

$$\text{high_crime}_{it} = \begin{cases} 1 & \text{if crime rate} > 0.04 \\ 0 & \text{otherwise} \end{cases}$$

- **Key regressors:**
 - `lprbarr` : log probability of arrest
 - `lprbconv` : log probability of conviction
 - `lpolpc` : log police per capita
 - `ldensity` : log population density
 - `lwcon` : log construction wage

Research question: What determines whether a county experiences high crime?

Linear Probability Model with FE

Question: How would you estimate a LPM with county fixed effects?

Model:

$$\text{high_crime}_{it} = \alpha_i + X'_{it}\beta + \varepsilon_{it}$$

where:

- α_i = county fixed effect
- β = **marginal effects** (direct interpretation)

Within transformation get rid off α_i :

1. Demeaning and get rid off α_i

$$\text{high_crime}_{it} - \text{high_crime}_i = (X_{it} - \bar{X}_i)' \beta + (\varepsilon_{it} - \bar{\varepsilon}_i)$$

2. Run FE regression over this transformed model

$$\text{high_crime}_{it} = X'_{it}\beta + \varepsilon_{it}$$

```
1 xtreg high_crime lprbarr lprbconv lpolpc ldensity lwcon, fe vce(cluster county)
```

LPM: The Problem

1. Predictions outside the $[0,1]$ range

$$\hat{P}(\text{high_crime} = 1) > 1 \quad (25 \text{ observations})$$

$$\hat{P}(\text{high_crime} = 1) < 0 \quad (180 \text{ observations})$$

2. Heteroskedasticity (by construction)

$$\text{Var}(\varepsilon_{it} \mid X_{it}) = P_{it}(1 - P_{it})$$

Depends on X_{it}

3. Constant marginal effects (unrealistic)

A 1% increase in arrest probability reduces high crime by **9.8 percentage points** for all counties, regardless of whether the baseline probability is 5% or 80%.

4. Imposed linearity

Nonlinear Model: Logit with FE

Model specification:

$$P(\text{high_crime} = 1 \mid X_{it}, \alpha_i) = \Lambda(X'_{it}\beta + \alpha_i)$$

where $\Lambda(z) = \frac{\exp(z)}{1+\exp(z)}$ is the **logistic CDF**

$$P(\text{high_crime}_{it} = 1 \mid X_{it}, \alpha_i) = \frac{\exp(\alpha_i + X'_{it}\beta)}{1 + \exp(\alpha_i + X'_{it}\beta)}$$

- Probabilities are **guaranteed** to be in $[0,1]$,
- Better fit for binary outcomes
- Non-constant marginal effects

The IPP problem

But estimating β and 90 α_i with only $T = 7$ observations creates the **Incidental Parameters Problem**

- As $N \rightarrow \infty$ but T fixed $\Rightarrow \hat{\beta}$ is **inconsistent**
- Bias = $O(1/T) \approx 14\%$

Solution: Sufficient Statistic

For the logit model,

$$S_i = \sum_{t=1}^T \text{high_crime}_{it}$$

is a **sufficient statistic** for α_i

Definition: S_i is sufficient for α_i if:

$$P(\text{high_crime}_{i1}, \dots, \text{high_crime}_{iT} \mid S_i, X_i, \beta) \text{ does not depend on } \alpha_i$$

What this means:

Instead of maximizing:

$$L(\beta, \alpha_1, \dots, \alpha_N \mid y, x) = \prod_{i=1}^N \prod_{t=1}^T P(y_{it} \mid \alpha_i, x_i, \beta)$$

Maximize the conditional likelihood:

$$L_c(\beta \mid y, x, S_1, \dots, S_N) = \prod_{i=1}^N P(y_i \mid S_i, x_i, \beta)$$

Result:

- α_i **disappears** from the likelihood
- Estimate β **without** estimating α_i
- $\hat{\beta}$ is **consistent** even with small T

Conditional Logit

In stata conditional Logit is implemented using `xtlogit` and `or` to interpret the coefficients as odds ratios (ORs):

```
1 xtlogit high_crime lprbarr lprbconv lpolpc ldensity lwcon, fe or
```

Results (N=98, 14 counties):

Variable	Odds Ratio	Interpretation
lprbarr	0.0002***	↓ Reduces odds 99.98%
lprbconv	0.0002***	↓ Reduces odds 99.98%
lpolpc	13,067***	↑ Increases odds (!)
ldensity	5,369	Not significant
lwcon	0.432	Not significant

Key findings:

- **Strong deterrence:** Higher arrest/conviction probability dramatically reduces crime odds
- **Police puzzle:** Positive coefficient likely reflects **reverse causality** (crime → hire police)
- **Sample note:** Only 14 “switcher” counties used (76 dropped due to no variation)

Interpretation: OR < 1 decreases odds, OR > 1 increases odds.

Nonparametric Density Estimation

Histograms

A nonparametric method to estimate the probability density function (PDF)

$$\hat{f}(x_0) = \frac{1}{N_h} \sum_{i=1}^N \frac{1}{2} \mathbf{1}\left(\left|\frac{x_i - x_0}{h}\right| < 1\right)$$

Idea: Count observations in bins, normalize by sample size and bin width

The degree of smoothing of the histogram can be controlled by selecting:

1. Number of bins

```
1 histogram price, bin(10) name(hist10, replace)
```

1. Bin width $2h$

```
1 histogram price, width(2000) density name(hist_wide, replace)
```

Kernel Density

Smooth, nonparametric estimate of the probability density function

$$\hat{f}(x_0) = \frac{1}{Nh} \sum_{i=1}^N K\left(\frac{x_i - x_0}{h}\right)$$

- Role of $K(\cdot)$: weighting function that determines how much each observation x_i contributes to the density estimate at the point x_0
- Role of h : smoothing parameter that controls the width of the kernel.

How to choose optimal h ?

- The one that minimizes MISE, but since we do not know f that is the true density function, we can assume that it follows a Normal distribution $f \approx N(\mu, \sigma^2)$.
- Silverman's Plug-in Solution is the one used by default in Stata.

$$\hat{h}^* = 1.364 \times \delta \times \min\left(s, \frac{IQR}{1.34}\right) \times N^{-1/5}$$

For Epanechnikov kernel:

$$\hat{h}^* = 0.9 \times \min\left(s, \frac{IQR}{1.34}\right) \times N^{-1/5}$$

Non Parametric Regression



Binned Scatter plots

With large N, scatter plots become unreadable, and can't see conditional means:

Binscatter Steps:

1. Divide tenure into 20 bins
2. Calculate mean of wage within each bin
3. Plot (mean tenure, mean wage) for each bin

```
1 scatter wage tenure
2 binscatter wage tenure, nq(20)
```




Controlling for Covariates

Frisch-Waugh-Lovell (FWL) Theorem:

1. Regress wage on controls $\rightarrow$ get residuals $w\tilde{age}$
2. Regress tenure on controls $\rightarrow$ get residuals $te\tilde{nure}$
3. Binscatter: plot $w\tilde{age}$ vs $te\tilde{nure}$

In Stata:


```
1 binscatter wage tenure, controls(age)
```



Binscatter by Groups

Compare relationships across subgroups:

```
1 binscatter wage tenure, nq(20) by(union)
```



Shows $E(wage|tenure)$ separately for union, useful for heterogeneity analysis.

Cattaneo et al. (2024)

Traditional binscatter issues:

1. Assumes linear CEF
2. **Bin choice matters** (10 vs 20 vs 50 bins)
3. **No formal inference** (confidence intervals)

binsreg improvements:

1. Employs semi-parametric methods to visualize CEF
2. **Data-driven bin selection** (optimal number of bins)
3. **Confidence intervals** (pointwise or uniform bands)

Nonparametric Regression

Kernel regression: Nadaraya-Watson Estimator

- `binscatter` can be seen as using an implicit rectangular (uniform) kernel.
- Kernel regression replaces this with a smooth weighting function to obtain a smooth approximation of the CEF:

$$\hat{m}_{NW}(x_0) = \frac{\sum_{i=1}^N K\left(\frac{x_i - x_0}{h}\right) y_i}{\sum_{i=1}^N K\left(\frac{x_i - x_0}{h}\right)}$$

- Points closer to x_0 get more weight
- Can be interpreted as weighted local average


Equivalent to:

$$\min_{\beta_0} \sum (y_i - \beta_0)^2 K\left(\frac{x_i - x_0}{h}\right)$$

Bandwidth and Bias–Variance Tradeoff

As we discussed, the election of h is which control the **variance-bias** trade off.

- Small h : very local averaging $\rightarrow$ low bias, high variance (overfitting)
- Large h : heavy smoothing $\rightarrow$ high bias, low variance (underfitting)
- Intermediate h : best compromise



Choosing optimal h :

Leave-One-Out Cross Validation

Toy data: $(x, y) = \{(1, 2), (2, 3), (3, 2), (4, 5)\}$

For a candidate bandwidth h :

1. Step 1: remove $(1, 2)$

Estimate $m_{1,h}^{\wedge}(1)$ using remaining three points.

2. Step 2: prediction error:

$$(2 - m_{1,h}^{\wedge}(1))^2$$

3. Step 3: Repeat for all points:

$$CV(h) = \frac{1}{4} \sum_{i=1}^4 (y_i - m_{i,h}^{\wedge}(x_i))^2$$

Choose h minimizing $CV(h)$.

In practice: `lpoly` uses plug-in bandwidth approximating this optimum.


Local Linear Regression

Instead of approximating the function close to x_0 with a constant, now we can approximate it with a line.

For each x_0 solves:

$$\min_{\beta_0, \beta_1} \sum (y_i - \beta_0 - \beta_1(x_i - x_0))^2 K\left(\frac{x_i - x_0}{h}\right)$$


$$m(\hat{x}_0) = \hat{\beta}_0, \quad m'(\hat{x}_0) = \hat{\beta}_1$$



npregress and Marginal Effects

```
1 npregress kernel wage tenure  
2 npgraph, title("Local Linear using npregress")
```

1. Allows multiple covariates
2. Allows to make inference using (SE, VCE, robust)
3. Allows to obtain marginal effects



Semiparametric regression models

Hedonic pricing equation of houses

- **Dataset:** HPRICE3 (Wooldridge)
- **Outcome variable:** Log of house prices

$$\ln \text{price}_i = X_i' \beta + \lambda(\ln \text{dist}_i) + u_i \quad \text{with} \quad E(u_i | X_i, \ln \text{dist}_i) = 0$$

- **Key regressors:**
 - **larea** : log of square footage of house
 - **lland** : log of square footage lot
 - **rooms** : # rooms in house
 - **bath** : # bathrooms
 - **age** : age of the house
 - **ldist** : log dist. from house to incin., feet

Research question: What was the effect of a local garbage incinerator on housing prices in North Andover in 1981?

House characteristics affect prices approximately **linearly**, but distance may have **nonlinear disamenity effects**.

Robinson Difference Estimator (PLM)

A nonparametric Frisch–Waugh “partialling-out” procedure.

1. Estimate conditional means (kernel):

$$E[\textit{lprice}|\textit{ldist}_i] = \hat{m}_y(\textit{ldist}_i) \quad \text{and} \quad E[X|\textit{ldist}_i] = \hat{m}_X(\textit{ldist}_i)$$

2. Partial out distance:

$$\textit{lp}\tilde{\textit{r}}\textit{ice}_i = \textit{lprice}_i - \hat{m}_y(\textit{ldist}_i) \quad \text{and} \quad \tilde{X}_i = X_i - \hat{m}_X(\textit{ldist}_i)$$

3. OLS:

$$\textit{lp}\tilde{\textit{r}}\textit{ice}_i = X_i^* \beta + u_i \quad \Rightarrow \quad \hat{\beta}_{PLM}$$

4. Recover nonlinear effect:

$$\hat{\lambda}(\textit{ldist}_i) = \hat{m}_y(\textit{ldist}_i) - \hat{m}_X(\textit{ldist}_i) \hat{\beta}$$

Idea: remove distance nonparametrically $\rightarrow$ estimate β with clean variation $\rightarrow$ reconstruct λ .

Implementation

```
1 semipar lprice larea lland rooms bath age if y81==1, ///
2     nonpar(ldist) robust ci ///
3     title("ldist non parametric") ///
4     xtitle("ldist") ///
5     ytitle("{&lambda}(ldist)")
```

Parametric part:

```
Number of obs =      142
R-squared      =    0.6863
Adj R-squared  =    0.6748
Root MSE      =    0.1859
```

	lprice	Coefficient	Std. err.	t	P> t	[95% conf. interval]	
	larea	.3266051	.0887425	3.68	0.000	.1511228	.5020874
	lland	.0790684	.0316665	2.50	0.014	.01645	.1416868
	rooms	.026588	.0234205	1.14	0.258	-.0197244	.0729003
	baths	.1611464	.0463598	3.48	0.001	.069473	.2528198
	age	-.0029953	.000969	-3.09	0.002	-.0049113	-.0010793

Implementation

```
1 semipar lprice larea lland rooms bath age if y81==1, ///
2   nonpar(ldist) robust ci ///
3   title("ldist non parametric") ///
4   xtitle("ldist") ///
5   ytitle("{&lambda}(ldist)")
```

Non parametric part:

